

# Introduction to Kubernetes

Dimitris Armenatzoglou  
Senior DevOps Engineer  
@ Smart Reporting

# Some linguistics first... What does “Kubernetes” means?

Koo-ber-ne-teez

From the Greek:

verb: **Κυβερνώ** - noun: **Κυβερνήτης**  
**To govern - governor**

Kee-ver-nee-tees

*The same etymological root for the words: **Cyber, cybernetics***

It also means:

*Captain, pilot or the helmsman of a ship*



# What is Kubernetes?

<https://kubernetes.io/docs/concepts/overview/what-is-kubernetes/>:

Kubernetes is a **portable, extensible**, open-source **platform** for managing **containerized workloads** and **services**, that facilitates both **declarative** configuration and **automation**.

<https://www.redhat.com/en/topics/containers/what-is-kubernetes>:

Kubernetes (also known as k8s or "kube") is an open source **container orchestration platform** that automates many of the manual processes involved in **deploying, managing**, and **scaling containerized applications**.

Why "K8S"?:

[https://en.wikipedia.org/wiki/Kubernetes#cite\\_note-4](https://en.wikipedia.org/wiki/Kubernetes#cite_note-4)

# Kubernetes Features

- Orchestrate containers on multiple hosts.
- Autoscaling
- Stateless and stateful applications
- Service discovery
- Extensible
- Rolling updates and Rollbacks
- Resource monitoring and Self-healing
- Storage and networking orchestration
- Native load-balancing
- ...

## **Most importantly:**


One common interface - the same API - for running containerized applications **across any cloud or any bare-metal environments.**

# Kubernetes Market Adoption


Kubernetes dominates the containers market.

According to the StackRox survey “The state of Container and Kubernetes Security Winter 2020” **86%** of the responders use Kubernetes in production.

Percentage of survey respondents using Kubernetes:



Which of the following container orchestrators do you use? (pick as many as apply)



# Kubernetes Architecture

## API-Server:

The interface to interact with the cluster. REST.

## Scheduler:

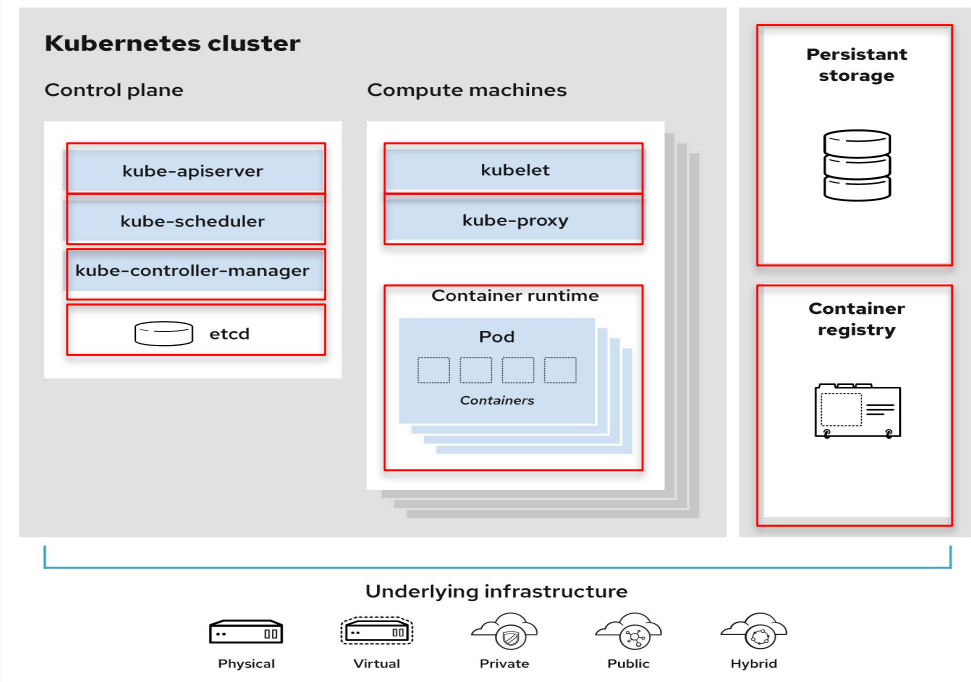
Schedules a POD to an appropriate Kubernetes node.

## Controller-manager:

Runs controller processes. E.g. Node-controller, replication controller, Endpoints controller etc.

## etcd:

Stores cluster state and configuration



## Kubelet:

It ensures that a set of PodSpecs is running healthy as containers on a node.

## Kube-Proxy:

A network proxy running on every node. Allows network communication to the Pods from inside or outside the cluster

## Container Runtime:

The engine which enables the cluster to run containers. E.g. Docker, rkt, containerd etc.

## Persistent Storage:

nfs, efs, ebs, glusterfs etc.

## Container Registry:

Harbor, Nexus, Docker Hub etc.

Some k8s key concepts:  
PODs, Nodes, Services,  
Controllers

# First things first... Enable Kubernetes on Docker Desktop

Install some k8s tools:

```
$ brew/yum/apt install  
kubectl
```

<https://kubernetes.io/docs/tasks/tools/install-kubectl/>

```
$ kubectl cluster-info
```

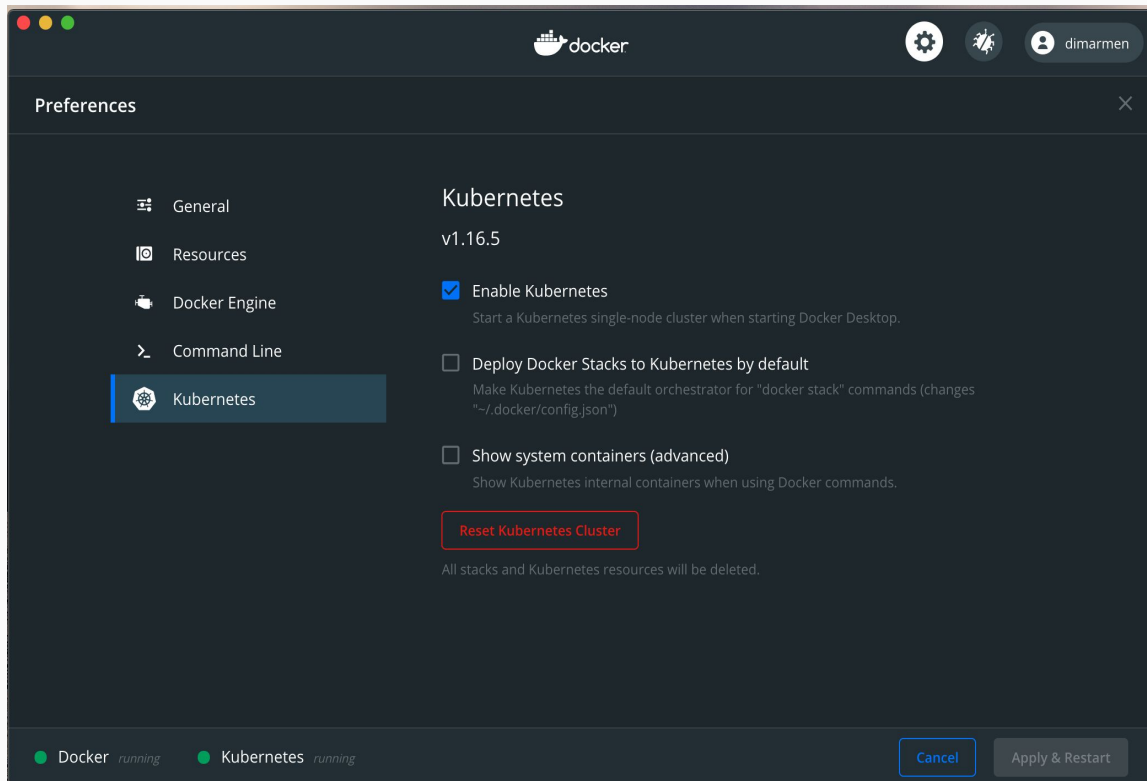
Workshop Manifests:

[https://github.com/DimArmen/kubernetes\\_demo](https://github.com/DimArmen/kubernetes_demo)

Some additional tools that you may find handy in the future:

<https://github.com/kubermatic/fubectl>  
<https://github.com/ahmetb/kubectx>

Kubernetes on Docker-Desktop:






# Pods

- The atomic unit of a kubernetes workload. The smallest deployable unit of work.
- A POD is **one or more** containers which form a cohesive instance of a running application.
- The PODS are **ephemeral**.
- Can we deploy Containers on Kubernetes?  
**Yes! But only inside Pods.**

```
apiVersion: v1
kind: Pod
metadata:
  name: web-1
  labels:
    app: web-1
spec:
  containers:
  - name: web
    image: nginx:1.18
    ports:
    - name: web
      containerPort: 80
      protocol: TCP
```

Pod name and labels that can be used by other controllers to identify this pod

The container name, image and ports

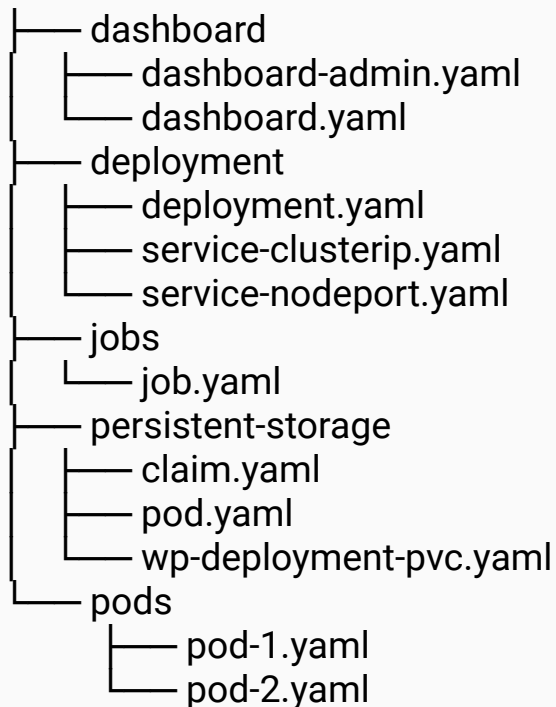


Kubernetes Cluster Node

<https://kubernetes.io/docs/reference/generated/kubernetes-api/v1.18/>

<https://kubernetes.io/docs/concepts/workloads/pods/>

# Create PODs



```
$ kubectl apply -f pods/pod-1.yaml
```

```
$ kubectl exec -it web-1 -- bin/sh
```

```
$ kubectl logs (-f) web-1
```

```
$ kubectl delete web-1
```

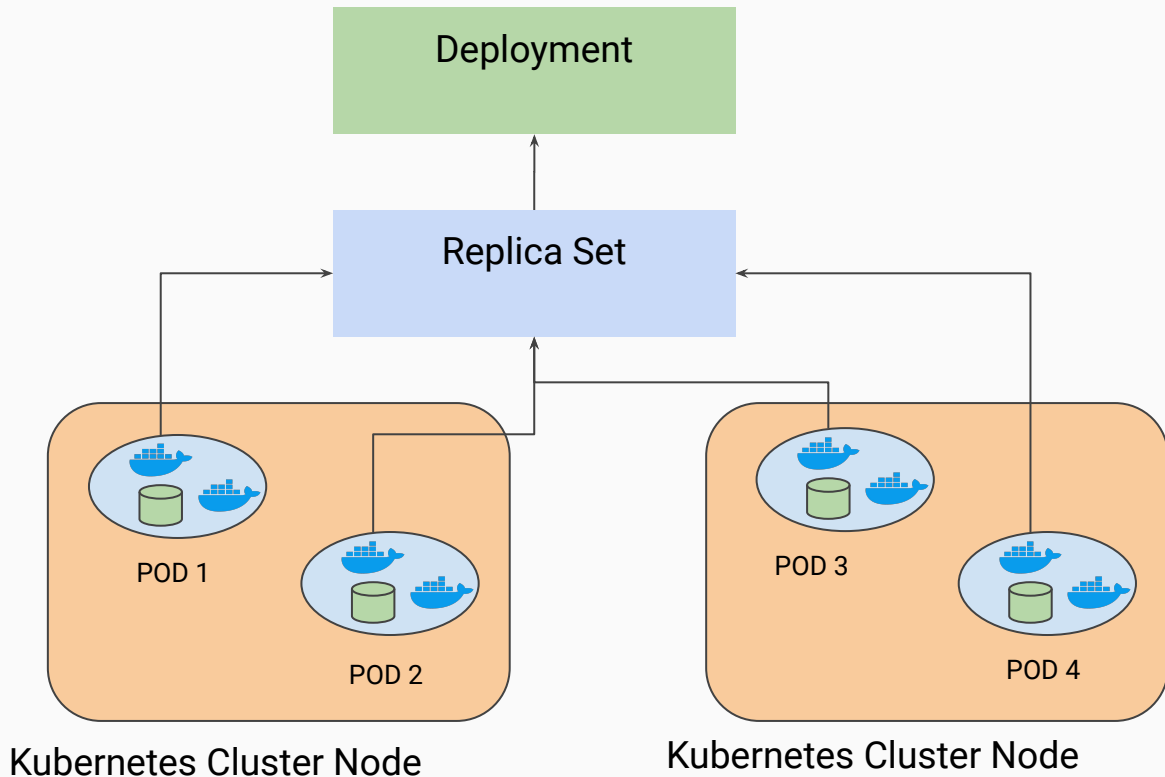
```
$ kubectl delete -f pods/pod-1.yaml
```

# Deployment

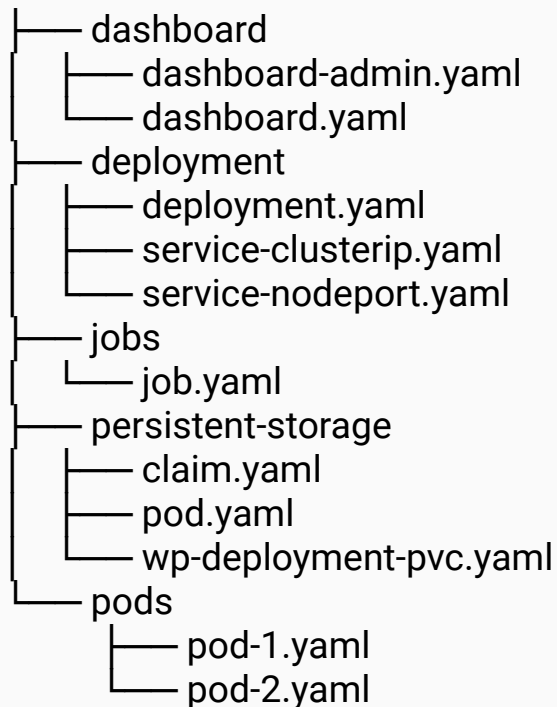
A Deployment Controller provides:

- Declarative updates for Pods via ReplicaSets.
- Update control and rollback functionality
- Uses POD templates to generate the replicaSets

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: deploy-example
spec:
  replicas: 3
  revisionHistoryLimit: 3
  selector:
    matchLabels:
      app: nginx
  strategy:
    type: RollingUpdate
    rollingUpdate:
      maxSurge: 1
      maxUnavailable: 1
  template:
    <pod template>
```



# Create Deployment



```
$ kubectl apply -f deployments/deployment.yaml
$ kubectl describe deployment hello-kubernetes
$ kubectl rollout history deployment hello-kubernetes
$ kubectl rollout history deployment hello-kubernetes --revision <number>
$ kubectl undo deployment hello-kubernetes
$ kubectl edit deployment hello-kubernetes
```

# Service

## Kubernetes Services:

- An endpoint used to expose an application running as a set of pods as a network service.
- Provide a reliable network endpoint for routing traffic to pods. **Not Ephemeral.**
- Autodiscovers associated PODS using matchLabels.
- Provides a DNS name such as:  
**http://<service\_name>:<service\_port>**

## Types of Services:

- ClusterIP
- NodePort
- LoadBalancer

```
apiVersion: v1
kind: Service
metadata:
  name: hello-kubernetes
spec:
  selector:
    name: hello-kubernetes
  type: ClusterIP | NodePort | LoadBalancer
  ports:
    - name: http
      port: 80
      targetPort: 8080
      nodePort: 32000
```

Pods Match labels


Service Port

Pods Port

If NodePort is used

## ClusterIP - Cluster Internal Only

- Cluster **internal only** communication.
- Finds the PODs by using the MatchLabels and groups them in a network service to forward traffic to them




# NodePort

A NodePort is exposing the Service on each Node.

From **outside the cluster** the Service is reachable at:


**<any\_node\_IP>:<NodePort>**



# LoadBalancer

Usually in a Cloud environment where Load Balancers are offered as a service (AWS ALB etc.)

The Kubernetes cloud-controller-manager will automatically create the load balancer resource.





# Job

Creates one or more pods and ensures that a specific number of them terminates successfully.

```
apiVersion: batch/v1
```

```
kind: Job
```

```
metadata:
```

```
  name: hello
```

```
spec:
```

```
  template:
```

```
    spec:
```

```
      containers:
```

```
      - name: hello
```

```
        image: busybox
```

```
        args: ["hello"]
```

```
        command: ["echo"]
```

```
      restartPolicy: Never
```

```
# completions: <int>
```

```
# backoffLimit: <int>
```

```
# parallelism: <int>
```

How many times should this job complete.

How many re-tries to attempt in case the job fails.

How many jobs can run in parallel.

# Persistent Volumes and Persistent Volume Claims


```
apiVersion: v1
kind: Pod
metadata:
  name: app
spec:
  containers:
  - name: app
    image: centos
    command: ["/bin/sh"]
    args: ["-c", "while true; do echo $(date
-u) >> /data/out.txt; sleep 5; done"]
    volumeMounts:
    - name: persistent-storage
      mountPath: /data
  volumes:
  - name: persistent-storage
    persistentVolumeClaim:
      claimName: pvc-example
```

```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: pvc-example
spec:
  accessModes:
  - ReadWriteOnce
  storageClassName: hostpath
  resources:
    requests:
      storage: 4Gi
```

Mount the volume specified by "name" at the "mountPath"

Use the PVC as a volume for this POD.

# Kubernetes Dashboard



## Deploy the Dashboard:

```
$ kubectl apply -f dashboard/dashboard.yaml
```

```
$ kubectl apply -f dashboard/dashboard-admin.yaml
```

```
$ kubectl -n kubernetes-dashboard describe secret $(kubectl -n kubernetes-dashboard get secret |  
grep admin-user | awk '{print $1}')
```

```
$ kubectl proxy
```

The Dashboard is available at:

[http://localhost:8001/api/v1/namespaces/kubernetes-dashboard/services/https:kubernetes-dashboard:/p  
roxy/](http://localhost:8001/api/v1/namespaces/kubernetes-dashboard/services/https:kubernetes-dashboard:/proxy/)

Thank you!